


```
void hello()
{
    std::cout << "Hello World! thread id = "
    << std::this_thread::get_id() << std::endl;
}
```

```
int main(int argc, char** argv)
{
    std::thread t(hello);
    return 0;
}
```

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$
```

```
clang++ -std=c++14 thread1.cpp
```

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ ./
```

```
a.out
```

```
libc++abi.dylib: terminating
```

```
Abort trap: 6
```

std::thread

- We ran `~thread()`; when we exited from the scope of `main` and it destroyed the thread object.
- **If `*this` has an associated thread `std::terminate()` is called**
- We hold a thread as long as:
`(joinable() == true)`
- We have to wait that the thread terminates its execution either calling `join` or `detach`

Fix

```
void hello()
{
    std::cout << "Hello World! thread id = " << std::this_thread::get_id() << std::endl;
}

int main(int argc, char** argv)
{
    std::thread t(hello);
    t.join();
    return 0;
}
```

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$
clang++ -std=c++14 thread1.cpp
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$
./a.out
Hello World! thread id = 0x102909000
```


will this code work?

```
int main(int argc, char** argv)
{
    // lambda function
    auto f = []()
    {
        std::cout << "ID of this thread = " << std::this_thread::get_id() << std::endl;
    };
    scoped_thread th(std::thread{f});
    return 0;
}
```

scoped thread


```
class scoped_thread
{
    std::thread t_;
public:
    explicit scoped_thread(std::thread t) : t_(std::move(t))
    {
        → if(t_.joinable() == false )
            std::logic_error("This is not a thread!!");

    }
    ~scoped_thread()
    {
        → if(t_.joinable())
            t_.join();
    }

    scoped_thread(scoped_thread&& x) : t_(std::move(x.t_))
    {}

    scoped_thread(scoped_thread&) = delete;
    scoped_thread& operator=(const scoped_thread&) = delete;
};
```


join vs detach



```
void complex_fnt( std::string& s)
{
    std::stringstream buff;
    std::ofstream f("./test.txt");
    //write to file string passed + thread id
    if(f.is_open())
    {
        buff << s << std::this_thread::get_id() << std::endl;
        f << buff.str();
        f.close();
    }
}
```

```
std::string s = "Hello world ";
```

```
//complex object passed by reference
std::thread t(complex_fnt, std::ref(s));
```

```
//assure job is done and return
t.join();
```

join vs detach

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ ./a.out 1
-- Join a complex task --
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ ls
a.out*          fancy_object.h  test.txt        thread2.cpp      thread4.cpp
async_check.cpp  scoped_thread.h thread1.cpp      thread3.cpp
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ cat test.txt
Hello world 0x103c27000
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$
```

join vs detach

```
void complex_fnt( std::string& s)
{
    std::stringstream buff;
    std::ofstream f("./test.txt");
    //write to file string passed + thread id
    if(f.is_open())
    {
        buff << s << std::this_thread::get_id() << std::endl;
        f << buff.str();
        f.close();
    }
}

std::string s = "Hello world ";
//complex object passed by reference
std::thread t(complex_fnt, std::ref(s));

//detach the thread... if the program exits before the threads
//completes its job, no job is done
t.detach();
```


join vs detach

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ ./a.out 2
-- Detach a complex task --
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ ls
a.out*          async_check.cpp  fancy_object.h  scoped_thread.h  thread1.cpp
thread2.cpp     thread3.cpp     thread4.cpp
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$
```

join vs detach

```
std::string s = "Hello world ";  
//complex object passed by reference  
std::thread t(complex_fnt, std::ref(s));  
  
//detach the thread... if the program exits before the threads  
completes its job, no job is done  
t.detach();  
  
//wait 5ms in order to assure that file is written on disk  
std::this_thread::sleep_for(std::chrono::nanoseconds(x));
```

join vs detach

```
./a.out 2 5000
-- Detach a complex task --
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ ls
a.out*          fancy_object.h  test.txt        thread2.cpp      thread4.cpp
async_check.cpp  scoped_thread.h thread1.cpp      thread3.cpp
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ cat test.txt
Hello world 0x10f069000
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$
```

Be sure all resources that your detached thread uses are still alive even when your program terminates

std::mutex and std::lock_guard

```
void fnt(int id, std::string s)
{
    std::cout << "id # = " << id
        << " - Functional object - I am a thread with ID = "
        << std::this_thread::get_id()
        << " custom msg = " << s
        << std::endl;
}
```

```
std::vector<std::thread> threads;
threads.emplace_back(fnt, 0, "Hi");
threads.emplace_back(fnt, 1, "Salut");
threads.emplace_back(fnt, 2, "Ciao");
threads.emplace_back(fnt, 3, "Hola");
```

```
for(auto& t: threads )
    t.join();
```

Output ... ???

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$  
clang++ -std=c++11 thread4.cpp  
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$  
./a.out 0
```

```
iiiidddd    ####    ====    0123    ----  
FFFFuuuunnnnccccttttiiiooonnnnaaaalll    oooobbbbjjjjeeecccctttt    ----  
IIII    aaammmm    aaaa    tttthhhrrrrreeeeaaaadddd    wwwiiiiitttthhhh  
IIIIDDDD    ====    0000xxx11110000777733440808147a000000000000  
ccccuuuussssttttoooommmm    mmmmssssgggg    ====    HSchiao  
laluoat
```

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$
```

std::mutex and std::lock_guard

```
std::mutex _m;

void fnt(int id, std::string s)
{
    std::lock_guard<std::mutex> lk{_m};

    std::cout << "id # = " << id
        << " - Functional object - I am a thread with ID = "
        << std::this_thread::get_id()
        << " custom msg = " << s
        <<std::endl;
}
```

Output

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ ./a.out 0
id # = 0 - Functional object - I am a thread with ID = 0x10051d000 custom msg = Hi
id # = 1 - Functional object - I am a thread with ID = 0x1005a0000 custom msg =
Salut
id # = 3 - Functional object - I am a thread with ID = 0x1006a6000 custom msg = Hola
id # = 2 - Functional object - I am a thread with ID = 0x100623000 custom msg = Ciao
```

for more details about mutexs
see: <http://en.cppreference.com/w/cpp/thread>

High level Interface

(not all APIs will be discussed)

Hardware Threads vs Software Threads

- 1 Core today == 2 hardware threads (due to hyper-threading)
- How can we understand how many software threads I can really run in parallel ???
- Usually more threads my application spawns (exceeding the number of hardware threads) and less job I get done.... let's see in few slides

Get physical threads

```
const unsigned physical_thread_number = std::thread::hardware_concurrency();  
  
std::cout << "Available physical thread = "  
<< available_threads << std::endl;
```

```
~/GitHub/cpp_sandbox/multithreading/thread_sample$ ./a.out 100 10
```

```
...  
Available physical thread = 4
```

```
...
```

std::async + std::future

```
std::mutex m;  
using lock = std::lock_guard<std::mutex>;  
std::map<std::thread::id, bool> ids;
```

```
void f(unsigned i)  
{
```

```
    lock lk{m};
```

```
    auto id = std::this_thread::get_id();
```

```
    std::cout << "thread #" << i << " id = " << id << std::endl;
```

```
    ids.insert(std::make_pair(id, false));
```

```
}
```

```
std::vector<std::future<void>> futures(10);
```

```
for(unsigned i = 0; i < 10; ++i)
```

```
    futures[i] = std::async(std::launch::any, f, i);
```

```
for(auto&f: futures )
```

```
    f.wait();
```

Possible output

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_spawn$ ./a.out
```

```
thread #0 id = 0x1041f0000
```

```
thread #1 id = 0x104273000
```

```
thread #2 id = 0x1042f6000
```

```
thread #3 id = 0x1041f0000
```

```
thread #4 id = 0x104273000
```

```
thread #5 id = 0x1042f6000
```

```
thread #6 id = 0x1041f0000
```

```
thread #7 id = 0x104273000
```

```
thread #8 id = 0x1042f6000
```

```
thread #9 id = 0x1041f0000
```

```
Actual thread spawned = 3
```

```
0x1041f0000
```

```
0x104273000
```

```
0x1042f6000
```

std::async

```
template< class Function, class... Args >  
std::future<typename std::result_of<Function(Args...)>::type>  
async( std::launch policy, Function&& f, Args&&... args );
```

- The template function `async` runs the function `f` asynchronously (potentially in a separate thread which may be part of a thread pool) and returns a `std::future` that will eventually hold the result of that function call.
- Policies to spawn computation are:
 - `std::launch::async`
 - `std::launch::deferred`
 - `launch::any (bitwise or async | deferred)`

std::future

```
template< class T > class future;  
template< class T > class future<T>;  
template<> class future<void>;
```


- The class template std::future provides a mechanism to access the result of asynchronous operations.
- You can use one of these objects to get a std::future back
 - std::promise
 - std::package_task
 - std::async

Write parallel code

Problem:

Compute parallel sum of an array of integers

N 10000
block size = 1250



```
template<typename I>
long async_accumulate(I begin, I end)
{
    const auto len = end - begin;
    → if (len <= block)
        return std::accumulate(begin, end, 0);


    I mid = begin + len / 2;
    → auto handle = std::async(async_accumulate<I>, mid, end);
    int sum = async_accumulate<I>(begin, mid);
    return sum + handle.get();
}
```

thread1

thread2

thread3

thread4



```
template<typename I>
long async_accumulate(I begin, I end)
{
    auto len = std::distance(begin, end);

    if (len == 0)
        return 0;

    const unsigned num_threads = compute_number_of_threads(len);
    std::vector<std::future<long>> res(num_threads-1);
    const unsigned block_size = len / num_threads;

    auto f = [](I b, I e, long init) { return std::accumulate(b,e,init); };

    auto start_block = begin;
    for(unsigned i = 0; i < num_threads-1; ++i)
    {
        auto end_block = start_block;
        std::advance(end_block, block_size);
        res[i] = std::async(f, start_block, end_block, 0);
        start_block = end_block;
    }

    long sum = f(start_block, end, 0);
    for(auto& f : res )
        sum += f.get();

    return sum;
}
```

Results

N = 10000 - Block = 1000

	Single thread	Divide/ Conquer	Array Split
	86us	614us	389us

N = 100000 - Block = 10000

	Single thread	Divide/ Conquer	Array Split
	510us	593us	458us


$N = 1000000 - \text{Block} = 100000$

	Single thread	Divide/ Conquer	Array Split
	~5.2 ms	~3.1 ms	~3.1 ms

Message Queue

Dining philosophers

Many producers
Many consumers



```

template<typename T>
class message_queue
{
    std::mutex _m;
    std::queue<T> _q;

    using lock_guard = std::lock_guard<std::mutex>;

public:
    using value_type = T;

    message_queue() = default;
    message_queue(const message_queue& ) = delete;
    message_queue& operator=(const message_queue&) = delete;

```

```

    void push(T item)
    {
        lock_guard lk{_m};
        _q.push(std::move(item));
    }

```

```

    bool pop(T& t)
    {
        lock_guard lk{_m};
        if(_q.empty())
            return false;

        t = std::move(_q.front());
        _q.pop();
        return true;
    }

```

Producer

```
void push(Q& q, unsigned id)
{
    for(unsigned i=id*LOOP_SIZE; i<(id+1)*LOOP_SIZE; ++i)
        q.push(i);
}
```

Consumer

```
void pop(Q& q)
{
    unsigned i;
    while(true)
    {
        q.pop(i);
    }
}
```

N Producers/ M Consumers

```
std::vector<std::future<void>> consumers(M);
std::vector<std::future<void>> producers(N);

int i = 0;
for (auto& p : producers)
    p = std::async(std::launch::async, push, std::ref(q), i++);

for (auto& c : consumers)
    c = std::async(std::launch::async, pop, std::ref(q));

for (auto& c : consumers)
    c.wait();

for (auto& p : producers)
    p.wait();
```


Signal events

Signal events

→ `std::condition_variable _cv;`

Push and notify

```
void push_and_notify(T item)
{
    {
        lock_guard lk{ _m };
        _q.push(std::move(item));
    }
    → _cv.notify_all();
}
```

sleep until there are new items to pop

```
void wait_and_pop(T& item)
{
    → std::unique_lock<std::mutex> lk{ _m };
    _cv.wait(lk, [this]{ return !_q.empty(); });
    item = std::move(_q.front());
    _q.pop();
}
```


- if you need more performances you can
 - decrease granularity of locks (example implementing a forward linked list with double pointers to head and tail)
 - lock free algorithms (more error prone though)

DeadLocks

```
std::vector<std::future<void>> tasks;

for(unsigned i = 0; i < 1; ++i)
{
    tasks.push_back(std::async(transfer,0,1,10));
    tasks.push_back(std::async(transfer,1,0,10));
}
```

```
void transfer(int from, int to, int sum)
{
```

```
    // deadlock
```

```
    auto& acc1 = _accounts[from];
```

```
    auto& acc2 = _accounts[to];
```



```
    lock_guard lk1(acc1.get_mutex());
```

```
    lock_guard lk2(acc2.get_mutex());
```

```
    std::cout << "Moving money from = " << from
<< " to = " << to << " sum = " << sum << "\n";
```

```
    if (acc1.balance() >= sum)
```

```
    {
```

```
        acc1.deposit(-sum);
```

```
        acc2.deposit(sum);
```

```
    }
```

```
}
```

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$  
clang++ -std=c++14 deadlock.cpp  
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$  
./a.out  
Moving money from = 0 to = 1 sum = 10
```

```
void transfer(int from, int to, int sum)
{
    // no deadlock
    auto& acc1 = _accounts[from];
    auto& acc2 = _accounts[to];

    → std::lock(acc1.get_mutex(), acc2.get_mutex());
    lock_guard lk1(acc1.get_mutex(), std::adopt_lock);
    lock_guard lk2(acc2.get_mutex(), std::adopt_lock);

    std::cout << "Moving money from = " << from
    << " to = " << to << " sum = " << sum << "\n";

    if (acc1.balance() >= sum)
    {
        acc1.deposit(-sum);
        acc2.deposit(sum);
    }
}
```

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$  
clang++ -std=c++14 deadlock.cpp  
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$  
./a.out  
Moving money from = 0 to = 1 sum = 10  
Moving money from = 1 to = 0 sum = 10  
Program end ... Balance 0 = 100 - Balance 1 = 0  
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$
```


Memory model

- The biggest introduction made in C++11
- Memory memory deals with:
 - atomic operations
 - visible effects of these operations
 - atomic access to object members given memory layout

is this code thread safe??

→

```
struct S
{
    long a;
    char b;
};
```

only C++11

```
//update a
auto f = [&s]()
{
    s.a = 10;
};

//upddate b
auto g = [&s]()
{
    s.b = 20;
};
```

→

```
//sample how to exploit memory model objects layout
auto fut1 = std::async(std::launch::async,f);
auto fut2 = std::async(std::launch::async,g);

fut1.wait();
fut2.wait();
```

std::atomic

```
template< class T > struct atomic
template<> struct atomic<Integral>;
template< class T > struct atomic<T*>;
```

- Each instantiation and full specialization of the std::atomic template defines an atomic type.
- Objects of atomic types are the only C++ objects that are free from data races; that is
- if one thread writes to an atomic object while another thread reads from it, the behaviour is well-defined.

```
std::atomic<double> b;  
std::atomic<long> c;  
std::atomic<int> d;  
std::atomic<short> e;  
std::atomic<char> f;  
std::atomic<long long> g;
```

```
std::cout << "double is atomic = "<< std::boolalpha << b.is_lock_free() << std::endl;  
std::cout << "long is atomic = "<< std::boolalpha << c.is_lock_free() << std::endl;  
std::cout << "int is atomic = "<< std::boolalpha << d.is_lock_free() << std::endl;  
std::cout << "short is atomic = "<< std::boolalpha << e.is_lock_free() << std::endl;  
std::cout << "char is atomic = "<< std::boolalpha << f.is_lock_free() << std::endl;  
std::cout << "long long is atomic = "<< std::boolalpha << g.is_lock_free() <<  
std::endl;
```

```
clang++ -std=c++14 atomic.cpp
```

```
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$ ./a.out
```

```
double is atomic = true
```

```
long is atomic = true
```

```
int is atomic = true
```

```
short is atomic = true
```

```
char is atomic = true
```

```
long long is atomic = true
```

```
std::atomic<int> ay{0};  
int x = 0;
```

```
void atomic_read_barrier()  
{  
→   std::cout << "y = " << ay.load() << std::endl;  
   std::cout << "x = " << x << std::endl;  
   std::cout << std::endl;  
}
```

```
void atomic_write_barrier()  
{  
→   x = 42;  
   ay.store(20);  
}
```

```
→ std::thread t2(atomic_read_barrier);  
std::thread t1(atomic_write_barrier);  
t1.join();  
t2.join();
```

x	y
0	0
42	0
42	20

```

void th_read_lock()
{
    lock lk{m};

    std::cout << "x = " << x << std::endl;
    std::cout << "y = " << y << std::endl;
    std::cout << std::endl;
}
void th_write_lock()
{
    lock lk{m};
    x = 42;
    y = 20;
}

```

```

std::thread t2(th_read_lock);
std::thread t1(th_write_lock);
t1.join();
t2.join();

```

x	y
0	0
42	20

Cool.. but can I craft
something useful with all
this stuff???


```
class spinlock
{
→ std::atomic_flag _flag;
public:
    spinlock() : _flag(ATOMIC_FLAG_INIT)
    {}

    spinlock(const spinlock&) = delete;
    spinlock& operator=(const spinlock&) = delete;
    spinlock(spinlock&&) = default;

    void lock()
    {
→ while(!_flag.test_and_set(std::memory_order_acquire));
    }

    void unlock()
    {
→ _flag.clear(std::memory_order_release);
    }
};
```

```

void th_write(spinlock& spin, int& a )
{
    spinlock_guard lk{spin};
    a = 42;
}

void th_read(spinlock& spin, int& a)
{
    spinlock_guard lk{spin};
    std::cout << " value = " << a << "\n";
}

```

```

nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$
clang++ -std=c++14 spinlock.cpp
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$ ./a.out
value = 42
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$ ./a.out
value = 42
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$ ./a.out
value = 0
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$ ./a.out
value = 0
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$ ./a.out
value = 42
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$ ./a.out
value = 42
nik@Nicolas-MacBook-Air:~/GitHub/cpp_sandbox/multithreading/thread_sync$ ./a.out
value = 42

```

Point to...

GitHub where I put all the code (there is also stuff I haven't shown)

https://github.com/nicola-cab/cpp_sandbox/tree/master/multithreading

Resources used

- <http://www.amazon.com/C-Concurrency-Action-Practical-Multithreading/dp/1933988770>
- <https://www.justsoftwaresolutions.co.uk/blog/>
- <http://herbsutter.com/category/effective-concurrency/>

Thanks...

Questions?